

Lab: Customized Embedded Processor Design for IoT

Tim Bücher, Stephan Holljes, Yannick Keller | February 6, 2019

CHAIR FOR EMBEDDED SYSTEMS



1 Introduction To The Project

- ADPCM
- Decoder
- Flow

2 Optimizations

- Performance
- Area


3 Benchmarks


4 Conclusion

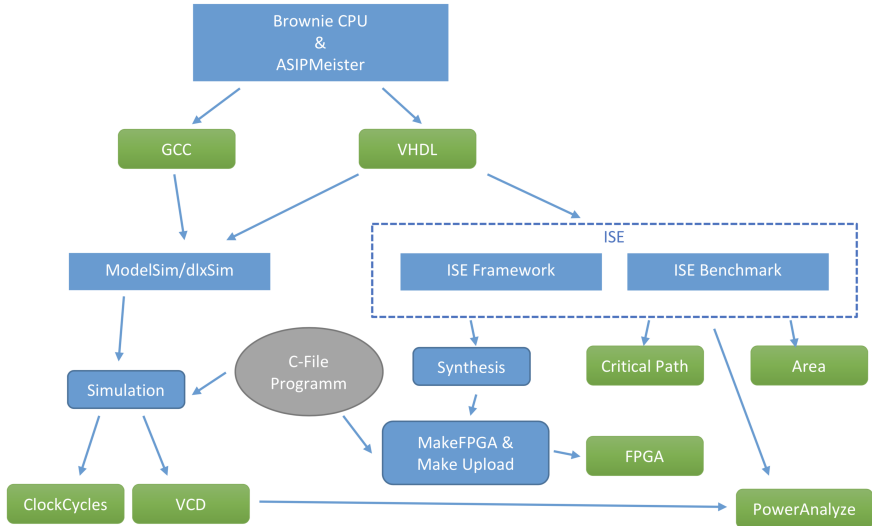
- Further work
- Challenges

Adaptive Differential Pulse Code Modulation (ADPCM)

- neighbouring audio samples often similar to each other
- variant of differential pulse-code modulation (DPCM)
- varies size of quantization step (adaptive)
- much less data to send/store







Performance and Area Optimization

Performance Optimization

- Optimized Clamping
- Optimized Value Prediction
- Added 3 extra ALUs

Performance Optimization

- Optimized Clamping
- Optimized Value Prediction
- Added 3 extra ALUs

Performance Optimization

- Optimized Clamping
- Optimized Value Prediction
- Added 3 extra ALUs

Performance optimization

```
114 /* Step 2 - Find new index value (for later) */
115 index += indexTable[delta];
116 if ( index < 0 ) index = 0;
117 if ( index > 88 ) index = 88;
118
119 /* Step 3 - Separate sign and magnitude */
120 sign = delta & 8;
121 delta = delta & 7;
122
123 /* Step 4 - Compute difference and new predicted value */
124 /*
125 ** Computes 'vpdiff = (delta+0.5)*step/4', but see comment
126 ** in adpcm_coder.
127 */
128 vpdiff = step >> 3;
129 if ( delta & 4 ) vpdiff += step;
130 if ( delta & 2 ) vpdiff += step>>1;
131 if ( delta & 1 ) vpdiff += step>>2;
132
133 if ( sign )
134     valpred -= vpdiff;
135 else
136     valpred += vpdiff;
137
138 /* Step 5 - clamp output value */
139 if ( valpred > 32767 )
140     valpred = 32767;
141 else if ( valpred < -32768 )
142     valpred = -32768;
143
144 /* Step 6 - Update step value */
145 step = stepsizeTable[index];
```

```
116 /* Step 2 - Find new index value (for later) */
117 index += indexTable[delta];
118
119 index = _builtin_brownie32_CLAMPI(index);
120
121 unsigned int arg = ((step<<16) & 0xFFFF0000) | delta;
122
123 valpred = _builtin_brownie32_VALPRED(arg, valpred);
124
125 valpred = _builtin_brownie32_CLAMPPV(valpred);
126
127 step = stepsizeTable[index];
```

```
19 lower_bound = "00000000000000000000000000000000";
20 upper_bound = "000000000000000000000000000001011000";
21
22
23
24 alu_cmp_alu_flag = ALU.cmp(source1, lower_bound);
25
26 // signed less than is either sign bit is 1 with no overflow
27 // or sign bit is 0 with overflow
28 alu_cmp_tmp_flag = alu_cmp_alu_flag[1:0];
29 alu_cmp_lt_cond1 = alu_cmp_tmp_flag == "10";
30 alu_cmp_lt_cond2 = alu_cmp_tmp_flag == "01";
31 clamp_low = alu_cmp_lt_cond1 | alu_cmp_lt_cond2;
32
33 alu_cmp_alu_flag2 = ALU1.cmp(upper_bound, source1);
34
35 // signed less than is either sign bit is 1 with no overflow
36 // or sign bit is 0 with overflow
37 alu_cmp_tmp_flag2 = alu_cmp_alu_flag2[1:0];
38 alu_cmp_lt_cond12 = alu_cmp_tmp_flag2 == "10";
39 alu_cmp_lt_cond22 = alu_cmp_tmp_flag2 == "01";
40 clamp_high = alu_cmp_lt_cond1 | alu_cmp_lt_cond2;
41
42 clamped = clamp_low | clamp_high;
43
44 clamped_value = (clamp_low) ? lower_bound:upper_bound;
45
46 result = (clamped) ? clamped_value:source1;
47
48 ForwardDataFromEXE(rd, result)
```

ASIP micro-operation description for clamping functions

```
33 delta = source1[3:0];
34 sign = delta[3];
35 zero = "00000000000000000000000000000000";
36 one = "00000000000000000000000000000001";
37 zero16 = "0000000000000000";
38 zero17 = "0000000000000000";
39 zero18 = "0000000000000000";
40 zero19 = "0000000000000000";
41 step_tmp1 = source1[31:16];
42 step_tmp2 = source1[31:17];
43 step_tmp3 = source1[31:18];
44 step_tmp4 = source1[31:19];
45 step = <zero16, step_tmp1>;
46 step_1s = <zero17, step_tmp2>;
47 step_2s = <zero18, step_tmp3>;
48 step_3s = <zero19, step_tmp4>;
49 delta3 = delta[2];
50 delta2 = delta[1];
51 delta1 = delta[0];
52
53
54 vpdiff_1 = (delta3) ? step:zero;
55 vpdiff_2 = (delta2) ? step_1s:zero;
56 vpdiff_3 = (delta1) ? step_2s:zero;
57
58 <temp1, flag1> = ALU.add(vpdiff_1, vpdiff_2);
59 <temp2, flag2> = ALU1.add(vpdiff_3, step_3s);
60 <temp3, flag3> = ALU2.add(temp1, temp2);
61 not_temp3 = ~temp3;
62 <temp4, flag4> = ALU3.add(one, not_temp3);
63
64 result = (sign) ? temp4:temp3;
65
66 ForwardDataFromEXE(rd, result)
```

ASIP micro-operation description for value prediction function



Area Optimization

- Removed all unused instructions (removed 20%)

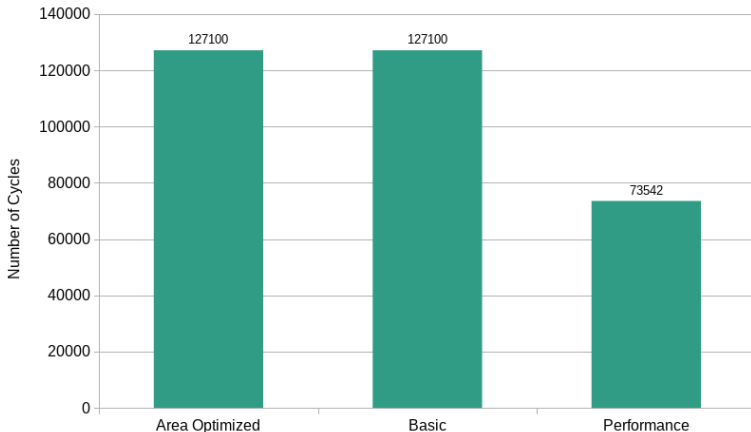
ADD	ELT	LW
SUB	ELTU	SB
MUL	EEQ	SH
DIV	ENEQ	SW
DIVU	ADDI	BRZ
MOD	SUBI	BRNZ
MODU	ANDI	JP
AND	ORI	JPL
NAND	XORI	TRAP
OR	LLSI	JPR
NOR	LRSI	JPRL
XOR	ARSI	NOP
LLS	LSOI	RETI
LRS	LB	EXBW
ARS	LH	EXHW

Available instructions


ADD	ELT	LW
SUB	ELTU	SB
MUL	EEQ	SH
DIV	ENEQ	SW
DIVU	ADDI	BRZ
MOD	SUBI	BRNZ
MODU	ANDI	JP
AND	ORI	JPL
NAND	XORI	TRAP
OR	LLSI	JPR
NOR	LRSI	JPRL
XOR	ARSI	NOP
LLS	LSOI	RETI
LRS	LB	EXBW
ARS	LH	EXHW

Used instructions


Benchmarks - Cycles




Benchmarks Area - LUTs



Basic
100%




Performance
+6%




Area
-12%


Benchmarks Area - Slices



Basic
100%




Performance
+16%



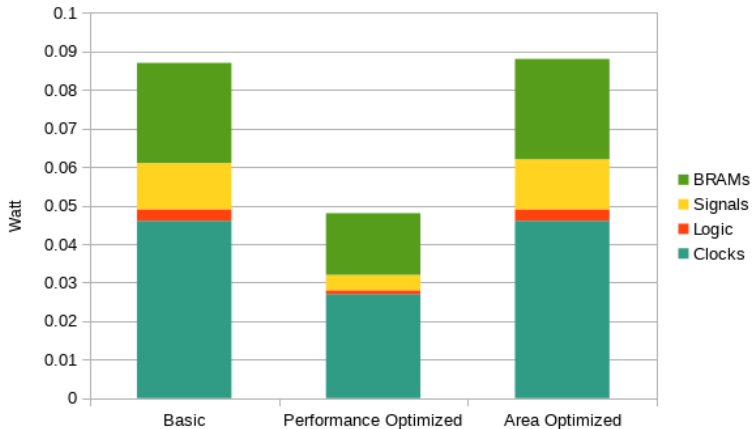
Area
-9%

Benchmarks - Power usage



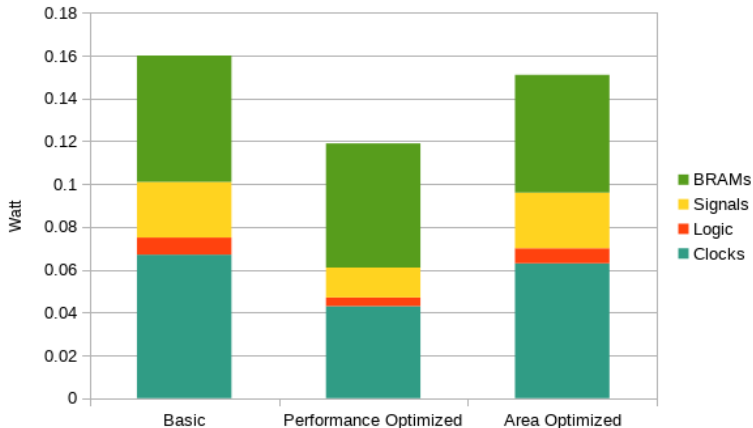
Power usage at 50MHz

Benchmarks - Power usage



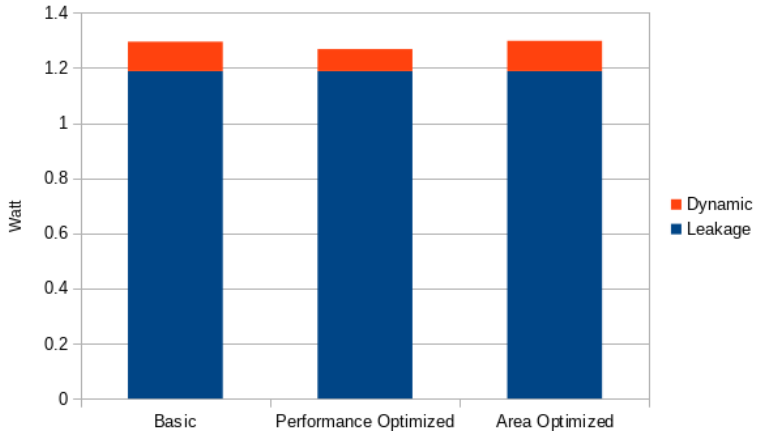
Power usage at lowest possible frequency (37MHz / 21.8MHz / 36MHz)

Benchmarks - Power usage




Power usage at highest possible frequency (84Mhz / 77MHz / 77Mhz)

Benchmarks - Leakage and Dynamic power




Dynamic and Leakage power at 50MHz

Benchmarks - Leakage and Dynamic power



Dynamic and Leakage power at lowest possible frequency (37MHz /
21.8MHz / 36MHz)

Benchmarks - Leakage and Dynamic power



Dynamic and Leakage power at highest possible frequency (84MHz / 77MHz / 77MHz)

Lessons learned

- Both optimizations are independent of each other
- Combination could reduce area and increase performance

- Decode more than one sample in one instruction
- Time spent waiting for memory write can calculate next sample

Reduce register size

- All values needed are limited to 16 bit
- Currently registers are 32 bit

Idea:

- Use a single instruction to decode a sample

Challenges:

- Too many ALUs and calculations increase critical path
- Data dependencies

Solution:

- Use multiple instructions for portions of the algorithm

Not enough ALUs

Idea:

- Use a single instruction to decode a sample

Challenges:

- Too many ALUs and calculations increase critical path
- Data dependencies

Solution:

- Use multiple instructions for portions of the algorithm

Idea:

- Use a single instruction to decode a sample

Challenges:

- Too many ALUs and calculations increase critical path
- Data dependencies

Solution:

- Use multiple instructions for portions of the algorithm

No easy way to add ROM

Idea:

- Use some sort of LUT for the ADPCM tables

Challenge:

- ASIPmeister does not allow to easily add ROM

Solution:

- Add a new Resource with custom VHDL code
- (Was not feasible for us in the time given)

No easy way to add ROM

Idea:

- Use some sort of LUT for the ADPCM tables

Challenge:

- ASIPmeister does not allow to easily add ROM

Solution:

- Add a new Resource with custom VHDL code
- (Was not feasible for us in the time given)

No easy way to add ROM

Idea:

- Use some sort of LUT for the ADPCM tables

Challenge:

- ASIPmeister does not allow to easily add ROM

Solution:

- Add a new Resource with custom VHDL code
- (Was not feasible for us in the time given)

Compiler too optimized (-fomit-instructions -fomg-fast)

- Compiling with gcc with -O1 -O2 and -O3 did not run on the FPGA
- However, it ran in Modelsim

Compiler too optimized (-fomit-instructions -fomg-fast)

- Compiling with gcc with -O1 -O2 and -O3 did not run on the FPGA
- However, it ran in Modelsim

CPU design too fragile?

Challenge:

- Removing unused instructions lead to failing execution on FPGA
- E.g. simply disabling the MUL instruction caused execution to fail before reaching main()

"Solution":

- Run in Modelsim

CPU design too fragile?

Challenge:

- Removing unused instructions lead to failing execution on FPGA
- E.g. simply disabling the MUL instruction caused execution to fail before reaching main()

"Solution":

- Run in Modelsim

CPU design too fragile?

Challenge:

- Removing unused instructions lead to failing execution on FPGA
- E.g. simply disabling the MUL instruction caused execution to fail before reaching main()

"Solution":

- Run in Modelsim

Challenge:

- High school grade reading comprehension
- Shortest paths in Xilinx PlanAhead are not the critical path in the CPU design
- MicroOp-Code syntax has documentation

"Solution":

- RTFM
- Re-read the timing reports and re-do benchmarks

Challenge:

- High school grade reading comprehension
- Shortest paths in Xilinx PlanAhead are not the critical path in the CPU design
- MicroOp-Code syntax has documentation

"Solution":

- RTFM
- Re-read the timing reports and re-do benchmarks

Challenge:

- High school grade reading comprehension
- Shortest paths in Xilinx PlanAhead are not the critical path in the CPU design
- MicroOp-Code syntax has documentation

"Solution":

- RTFM
- Re-read the timing reports and re-do benchmarks

Challenge:

- High school grade reading comprehension
- Shortest paths in Xilinx PlanAhead are not the critical path in the CPU design
- MicroOp-Code syntax has documentation

"Solution":

- RTFM
- Re-read the timing reports and re-do benchmarks

Finding things to read is hard too

Challenge:

- ASIPmeister is somewhat old, documentation hard to find (or in Japanese)
- Searching the web for browniestd32 returns mostly recipes
- (Minor) Versions in the docs don't match used versions

"Solution":

- Trial and error

Finding things to read is hard too

Challenge:

- ASIPmeister is somewhat old, documentation hard to find (or in Japanese)
- Searching the web for browniestd32 returns mostly recipes
- (Minor) Versions in the docs don't match used versions

"Solution":

- Trial and error

Challenge:

- In general a longer turn-around time for feedback than in usual software development
- "Blind coding" (no linting or other tools, feedback during compilation)
- Code generation (VHDL files, binutils) takes especially long
- Debugging the FPGA was unfeasible for us
- (For one day ASIPmeister wouldn't start at all, complaining about license issues)

Solution:

- Patience

Challenge:

- In general a longer turn-around time for feedback than in usual software development
- "Blind coding" (no linting or other tools, feedback during compilation)
- Code generation (VHDL files, binutils) takes especially long
- Debugging the FPGA was unfeasible for us
- (For one day ASIPmeister wouldn't start at all, complaining about license issues)

Solution:

- Patience